


Linear Regression

By Dr. Nzamba Bignoumba



$$y = -ax + b$$

Average training duration: 4 hours 00 minute

Outline

Machine learning overview	→	15 min
Linear regression: theory	→	45 min
Linear regression: use case	→	01.30 h
Model deployment	→	01.30 h

Machine learning overview

AI

what can we do with
it?

Natural Language Processing
Computer Vision
Signal Processing
⋮

Machine learning overview

Content summary of one or more documents
Translation from one language to another
Code generation
⋮

OpenAI-ChatGPT
DeepSeek
GitHub Copilot
Cursor | Codex

Visual content generation
Medical image classification
Agricultural Image Classification
Object detection
⋮

Midjourney
Canva
Aidoc
IA Agri
AgriHyphen AI

Natural Language Processing
Computer Vision
Signal Processing
⋮

Weather forecast
Disease and mortality forecasts/predictions
Stock market forecasts
Electricity consumption forecasts
Anomaly detection (cybersecurity)
⋮

AWS SageMaker - DeepAR
Nixtla-TimeGPT
Meta-Prophet
Zindi Africa
Amini

Machine learning overview

AI

what can we do with
it?



Natural Language
Processing Computer Vision
Signal Processing
⋮

With what type of
algorithms?



Machine Learning
Deep Learning

Linear regression

Machine learning overview

K-NN

Support Vector
Machines

Logistic
Regression

K-Means

Decision Trees

Gradient Boosting Machines

Random Forest

Linear
Regression

Principal Component
Analysis

Machine Learning

...

Deep Learning

Generative Adversarial
Neural Network

Varational
Autoencoder

Feed Forward
Neural Network

State Space Model

Word
Embeddings

Autoencoder

Graph Neural Network

Neural Ordinary
Differential Equations

Diffusion Model

Recurrent
Neural Network

Normalizing Flows

Transformer

Neural Radiance Field

...

Machine learning overview

IA

what can we do with
it?

Natural Language
Processing Computer Vision
Signal Processing
⋮

With what type of
algorithms?

Machine Learning

Deep Learning

How do they work ?

Supervised
Non-Supervised
Self-Supervised
⋮

Machine learning overview

Let's have the dataset \mathbf{X} s and its corresponding target \mathbf{y} s. The model will use \mathbf{X} to predict $\hat{\mathbf{y}}$, $\text{model}(\mathbf{X}) = \hat{\mathbf{y}}$. Then, $\hat{\mathbf{y}}$ will be compared to \mathbf{y} to calculate the error \mathbf{e} made by the model, $|\mathbf{y} - \hat{\mathbf{y}}| = \mathbf{e}$. This error serves to **refine the model parameters** so that it can predict a value $\hat{\mathbf{y}}$ very close to \mathbf{y} .

Observation

$\mathbf{x} \rightarrow \mathbf{y}$

Target

\mathbf{e}

The error that the model
seeks to minimize.

Model

$\hat{\mathbf{y}}$

Prediction



- Supervised
- Non-Supervised
- Self-Supervised
- ⋮

Machine learning overview

We only have **Xs** data. No matching **y** targets are available. The model will leverage data **similarities** and **co-occurrences** to perform the assigned task. For example, the clustering task, which consists of grouping data with similar patterns.

Observation

Unavailable

X ~~**y**~~

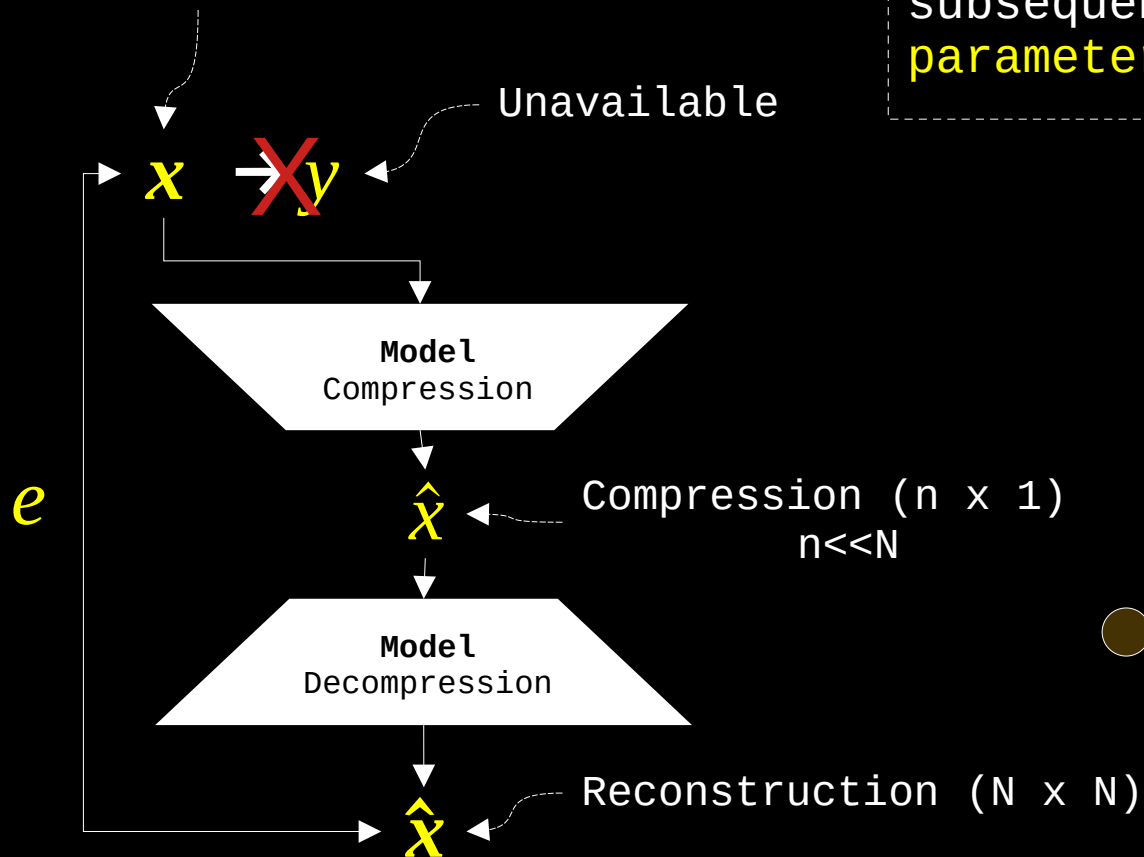
Model

Clustering

Supervised
Non-Supervised
Self-Supervised
⋮

Machine learning
overview

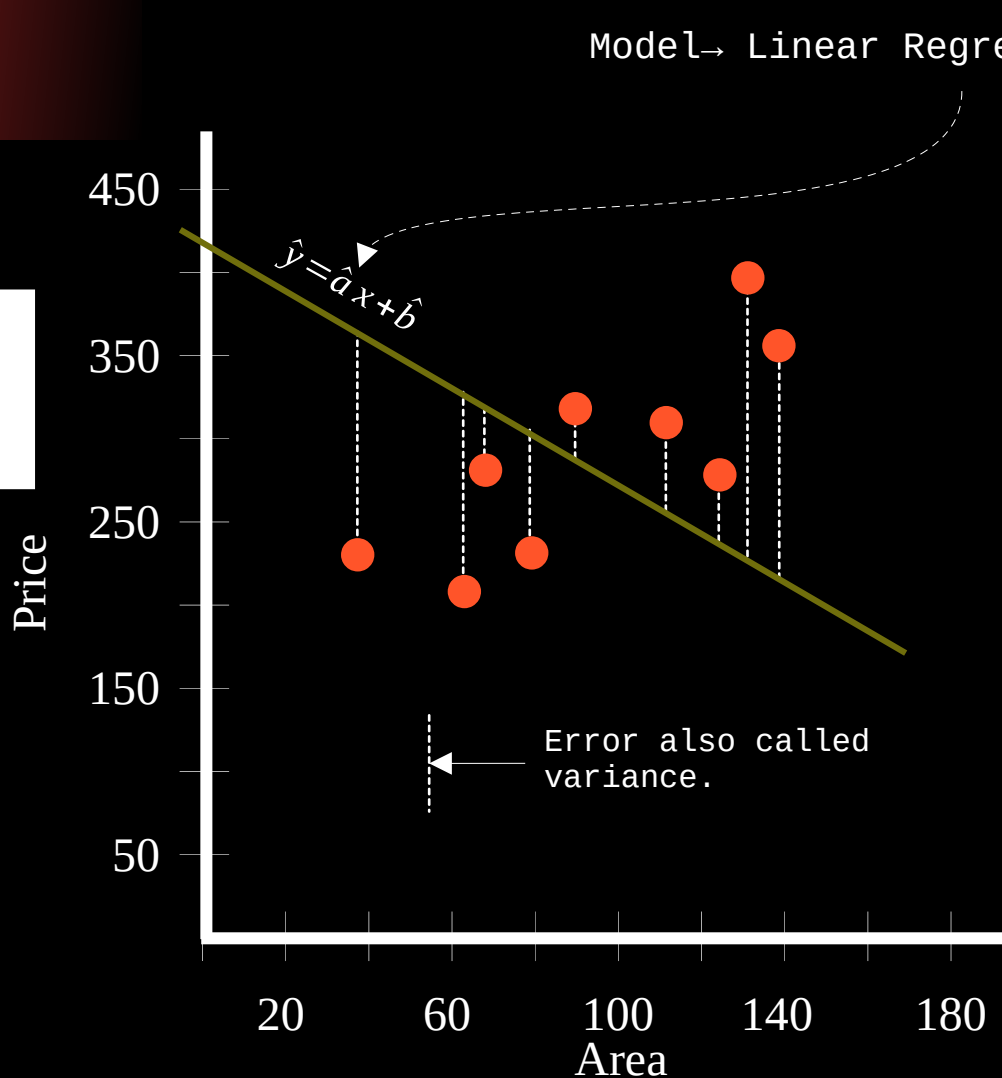
We only have the training dataset \mathbf{Xs} . The model takes an observation \mathbf{X} as input, $\text{model}(\mathbf{X})=\hat{\mathbf{X}}$, and compares its prediction $\hat{\mathbf{X}}$ to this same observation \mathbf{X} . The error $|\mathbf{X}-\hat{\mathbf{X}}|=\mathbf{e}$, will subsequently be used to fit the model parameters.

Observation ($N \times N$)

Supervised
Non-Supervised
Self-Supervised
⋮

Theory

Price chart (in XAF)
of house locations
based on their
surface area (in m²).



The training consists of finding the values \hat{a} and \hat{b} that minimize the error (E) as much as possible.


$$E = | \quad | + | \quad | + \dots + | \quad |$$

$$E = (e_1 + e_2 + \dots + e_9) / I$$

$I=9$, number of data points.


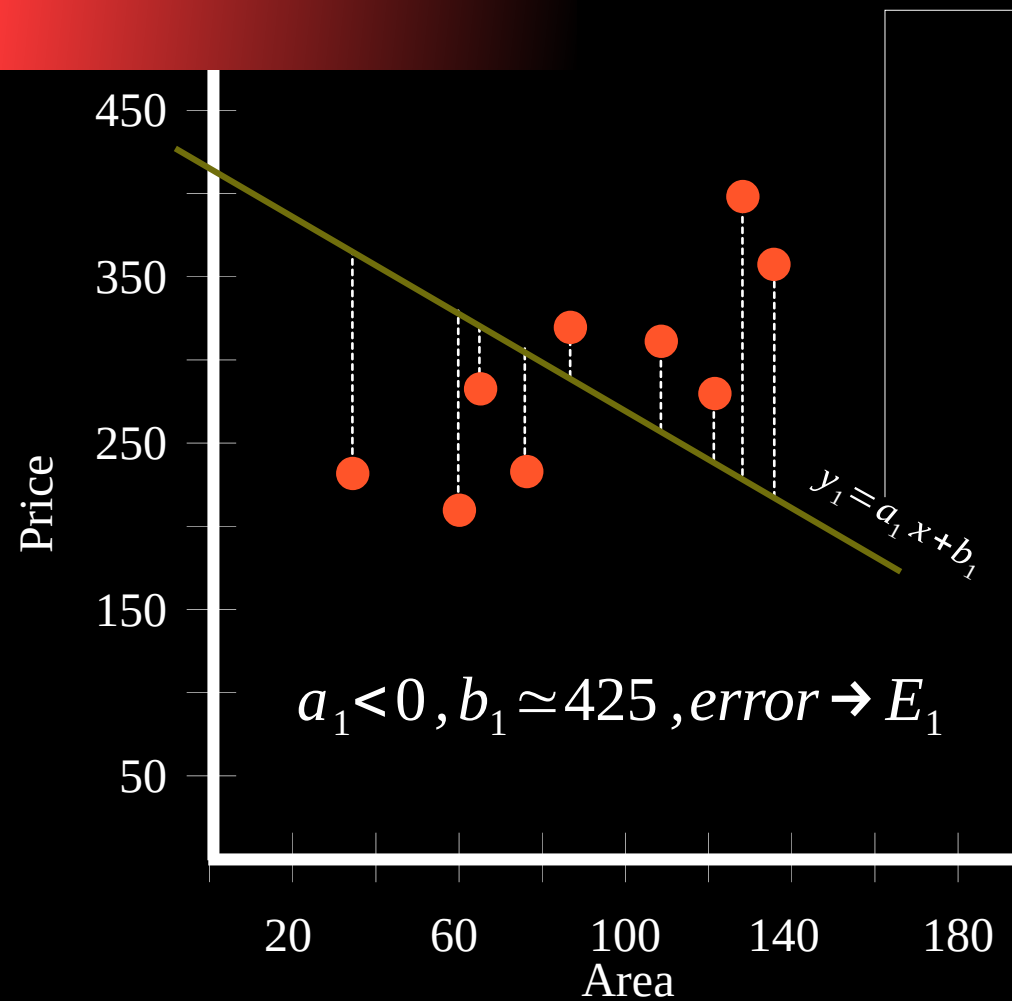
Theory

Price chart (in XAF)
of house locations
based on their
surface area (in m²).



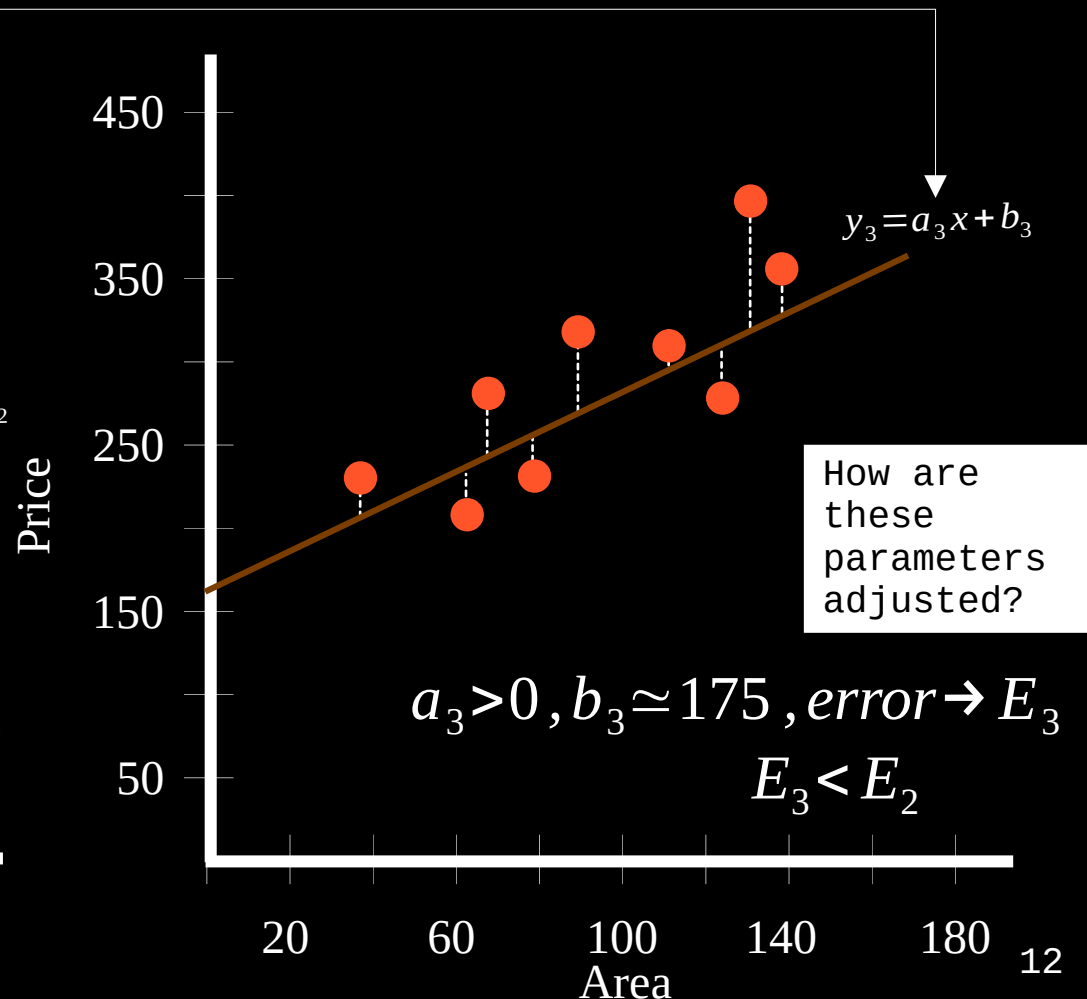
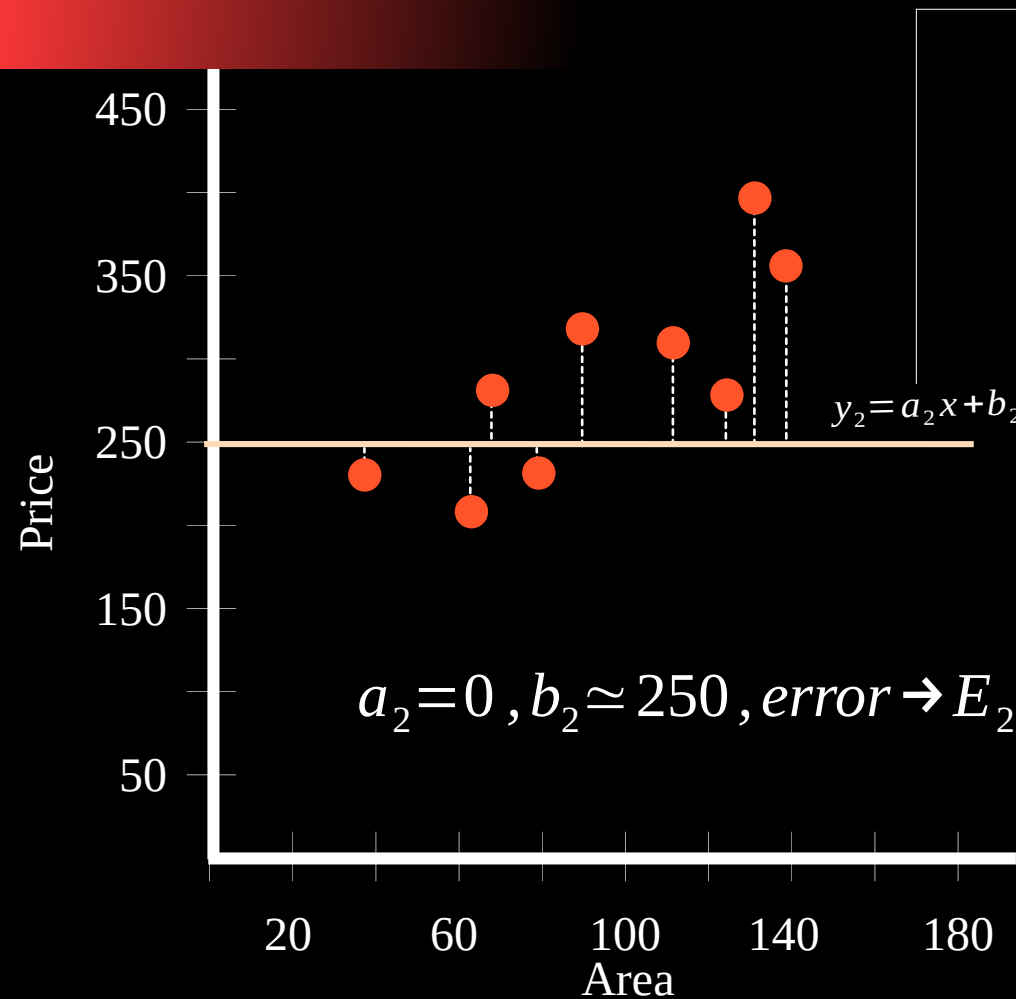
Theory

Model improvement after n iterations

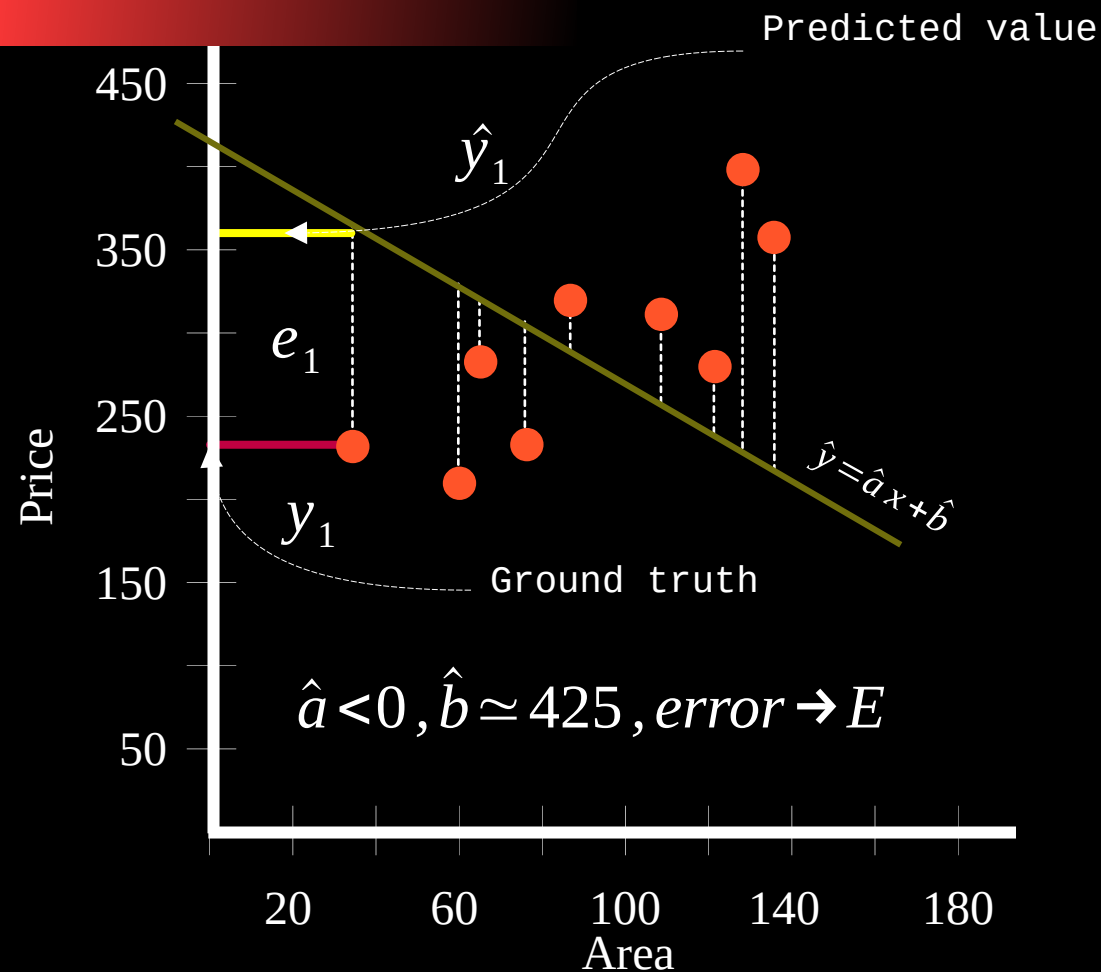


Theory

Model improvement after $n+k$ ($k>0$) iterations



Theory



$$e_1 = (y_1 - \hat{y}_1)^2 = [y_1 - (\hat{a}x_1 + \hat{b})]^2$$

$$E = (e_1 + e_2 + \dots + e_9) / I = \frac{1}{I} \sum_{i=1}^I e_i$$

E is called **Mean Squared Error (MSE)** and I=9, is the number of data points.

$$E = \frac{1}{I} \sum_{i=1}^I [y_i - (\hat{a}x_i + \hat{b})]^2$$

How to minimize error E?

By finding the optimal values of \hat{a} and \hat{b} that minimize this error.

Let's differentiate E to find the optimal parameters.

Theory

$$E = \frac{1}{I} \sum_{i=1}^I [y_i - (\hat{a} x_i + \hat{b})]^2$$

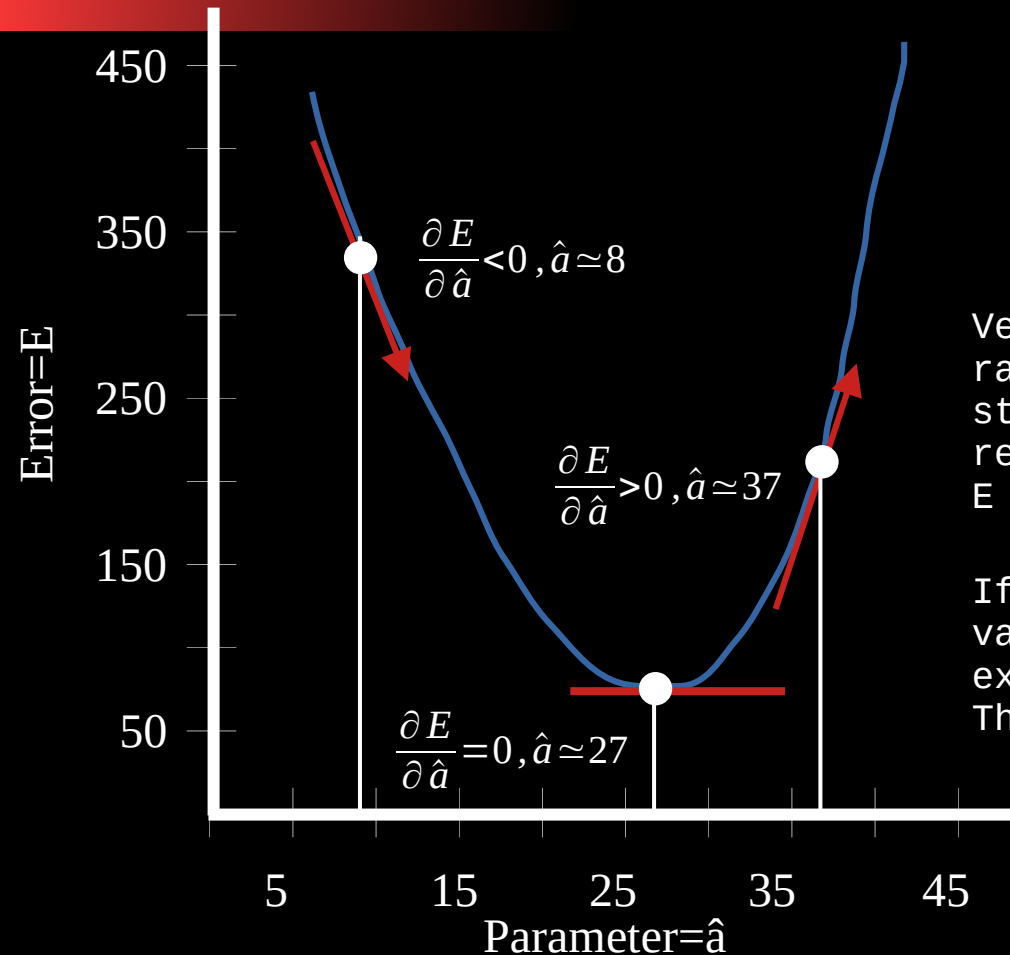
$$[(u - v)^2]' = 2(u - v)(u' - v'), u = y_i \text{ and } v = (\hat{a} x_i + \hat{b})$$

$$\frac{\partial E}{\partial \hat{a}} = \frac{1}{I} \sum_{i=1}^I 2[y_i - (\hat{a} x_i + \hat{b})](-x_i) = -\frac{2}{I} \sum_{i=1}^I [y_i - (\hat{a} x_i + \hat{b})](x_i)$$

$$\frac{\partial E}{\partial \hat{b}} = \frac{1}{I} \sum_{i=1}^I 2[y_i - (\hat{a} x_i + \hat{b})](-1) = -\frac{2}{I} \sum_{i=1}^I [y_i - (\hat{a} x_i + \hat{b})]$$

Theory

Reminder on the derivatives of functions



Suppose we want to calculate the derivative of E at the point $\hat{a}=x$, formally we have:

$$E'(\hat{a}=x) = \frac{\partial E}{\partial \hat{a}}$$

Verbally, the above function means: what is the rate of change of E (decrease, increase, stagnation), when we are at point \hat{a} . This rate is represented by a vector tangent to the curve of E at \hat{a} .

If we want to minimize E , we need to find the value \hat{a} that leads to this minimization. In our example, this value is approximately equal to 27. The same process is carried out for \hat{b} .

Let's calculate the derivative of E as a function of \hat{a} and \hat{b} .

Theory

Linear algebra reminder

Addition of
vectors

$$\begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_k \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_k \end{bmatrix} = \begin{bmatrix} a_1 + b_1 \\ a_2 + b_2 \\ \vdots \\ a_k + b_k \end{bmatrix}$$

Exemple $\rightarrow \begin{bmatrix} 2.5 \\ 18.0 \\ \vdots \\ 7.5 \end{bmatrix} + \begin{bmatrix} 3.9 \\ 8.1 \\ \vdots \\ 4.2 \end{bmatrix} = \begin{bmatrix} 6.4 \\ 26.1 \\ \vdots \\ 11.7 \end{bmatrix}$

Product of a
scalar with a
vector

$$X \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_k \end{bmatrix} = \begin{bmatrix} Xa_1 \\ Xa_2 \\ \vdots \\ Xa_k \end{bmatrix}$$

Exemple $\rightarrow 2 \begin{bmatrix} 3.9 \\ 8.1 \\ \vdots \\ 4.2 \end{bmatrix} = \begin{bmatrix} 7.8 \\ 16.2 \\ \vdots \\ 8.4 \end{bmatrix}$

Theory

Linear algebra reminder

To multiply two matrices, the number of columns in the first matrix must be equal to the number of rows in the second.

Product of matrices

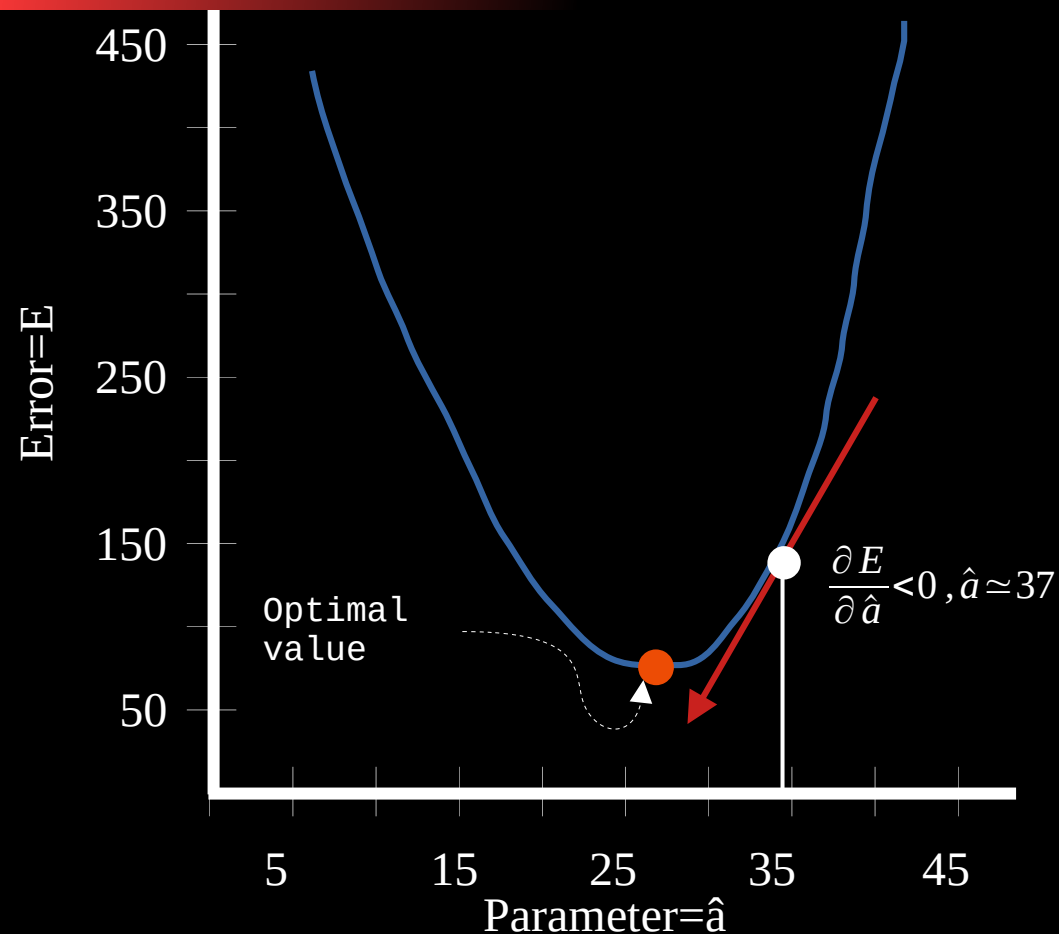
$$\begin{matrix} 2 \times 2 & 2 \times 3 & & 2 \times 3 \\ \begin{bmatrix} a_1^1 & a_2^1 \\ a_1^2 & a_2^2 \end{bmatrix} & \begin{bmatrix} b_1^1 & b_2^1 & b_3^1 \\ b_1^2 & b_2^2 & b_3^2 \end{bmatrix} & = & \begin{bmatrix} a_1^1 b_1^1 + a_2^1 b_1^2 & a_1^1 b_2^1 + a_2^1 b_2^2 & a_1^1 b_3^1 + a_2^1 b_3^2 \\ a_1^2 b_1^1 + a_2^2 b_1^2 & a_1^2 b_2^1 + a_2^2 b_2^2 & a_1^2 b_3^1 + a_2^2 b_3^2 \end{bmatrix}
 \end{matrix}$$

Exemple

$$\begin{bmatrix} 2,1 & 4,3 \\ 1,7 & 7,0 \end{bmatrix} \begin{bmatrix} 2,5 & 10,0 & 13,2 \\ 5,4 & 4,4 & 6,5 \end{bmatrix} = \begin{bmatrix} 2,1*2,5+4,3*5,4 & 2,1*10,0+4,3*4,4 & 2,1*13,2+4,3*6,5 \\ 1,7*2,5+7,0*5,4 & 1,7*10,0+7,0*4,4 & 1,7*13,2+7,0*6,5 \end{bmatrix}$$

$$\begin{bmatrix} 2,1 & 4,3 \\ 1,7 & 7,0 \end{bmatrix} \begin{bmatrix} 2,5 & 10,0 & 13,2 \\ 5,4 & 4,4 & 6,5 \end{bmatrix} = \begin{bmatrix} 28,47 & 39,92 & 55,67 \\ 42,05 & 47,8 & 67,94 \end{bmatrix}$$

Theory



The new value of \hat{a} will be:

$$\hat{a} = \hat{a} - \eta \frac{\partial E}{\partial \hat{a}}$$

$\eta > 0$ Is the **learning rate**. It is often set at 0.001.

This adjustment process called **gradient descent** is repeated n times until the optimal value of \hat{a} (as well as \hat{b}) that minimizes E is found.

What happens if we have several features?

Theory

$$\hat{y} = \hat{a}_1 x_1 + \hat{a}_2 x_2 + \dots + \hat{a}_k x_k + \hat{b} = [x_1, x_2, \dots, x_k] \begin{bmatrix} \hat{a}_1 \\ \hat{a}_2 \\ \vdots \\ \hat{a}_k \end{bmatrix} + \hat{b}$$

$$\hat{y} = X\theta + \hat{b} \mid X = [x_1, x_2, \dots, x_k], \theta = \begin{bmatrix} \hat{a}_1 \\ \hat{a}_2 \\ \vdots \\ \hat{a}_k \end{bmatrix}$$

$$E = \frac{1}{I} \sum_{i=1}^I [y_i - (X_i \theta + \hat{b})]^2$$

$$X = [x_1, x_2, \dots, x_k] \rightarrow X^T = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \end{bmatrix}$$

$$\nabla_{\theta} E(\theta) = \frac{1}{I} \sum_{i=1}^I 2[y_i - (X_i \theta + \hat{b})](-X_i)^T = -\frac{2}{I} \sum_{i=1}^I [y_i - (X_i \theta + \hat{b})](X_i)^T$$

Theory

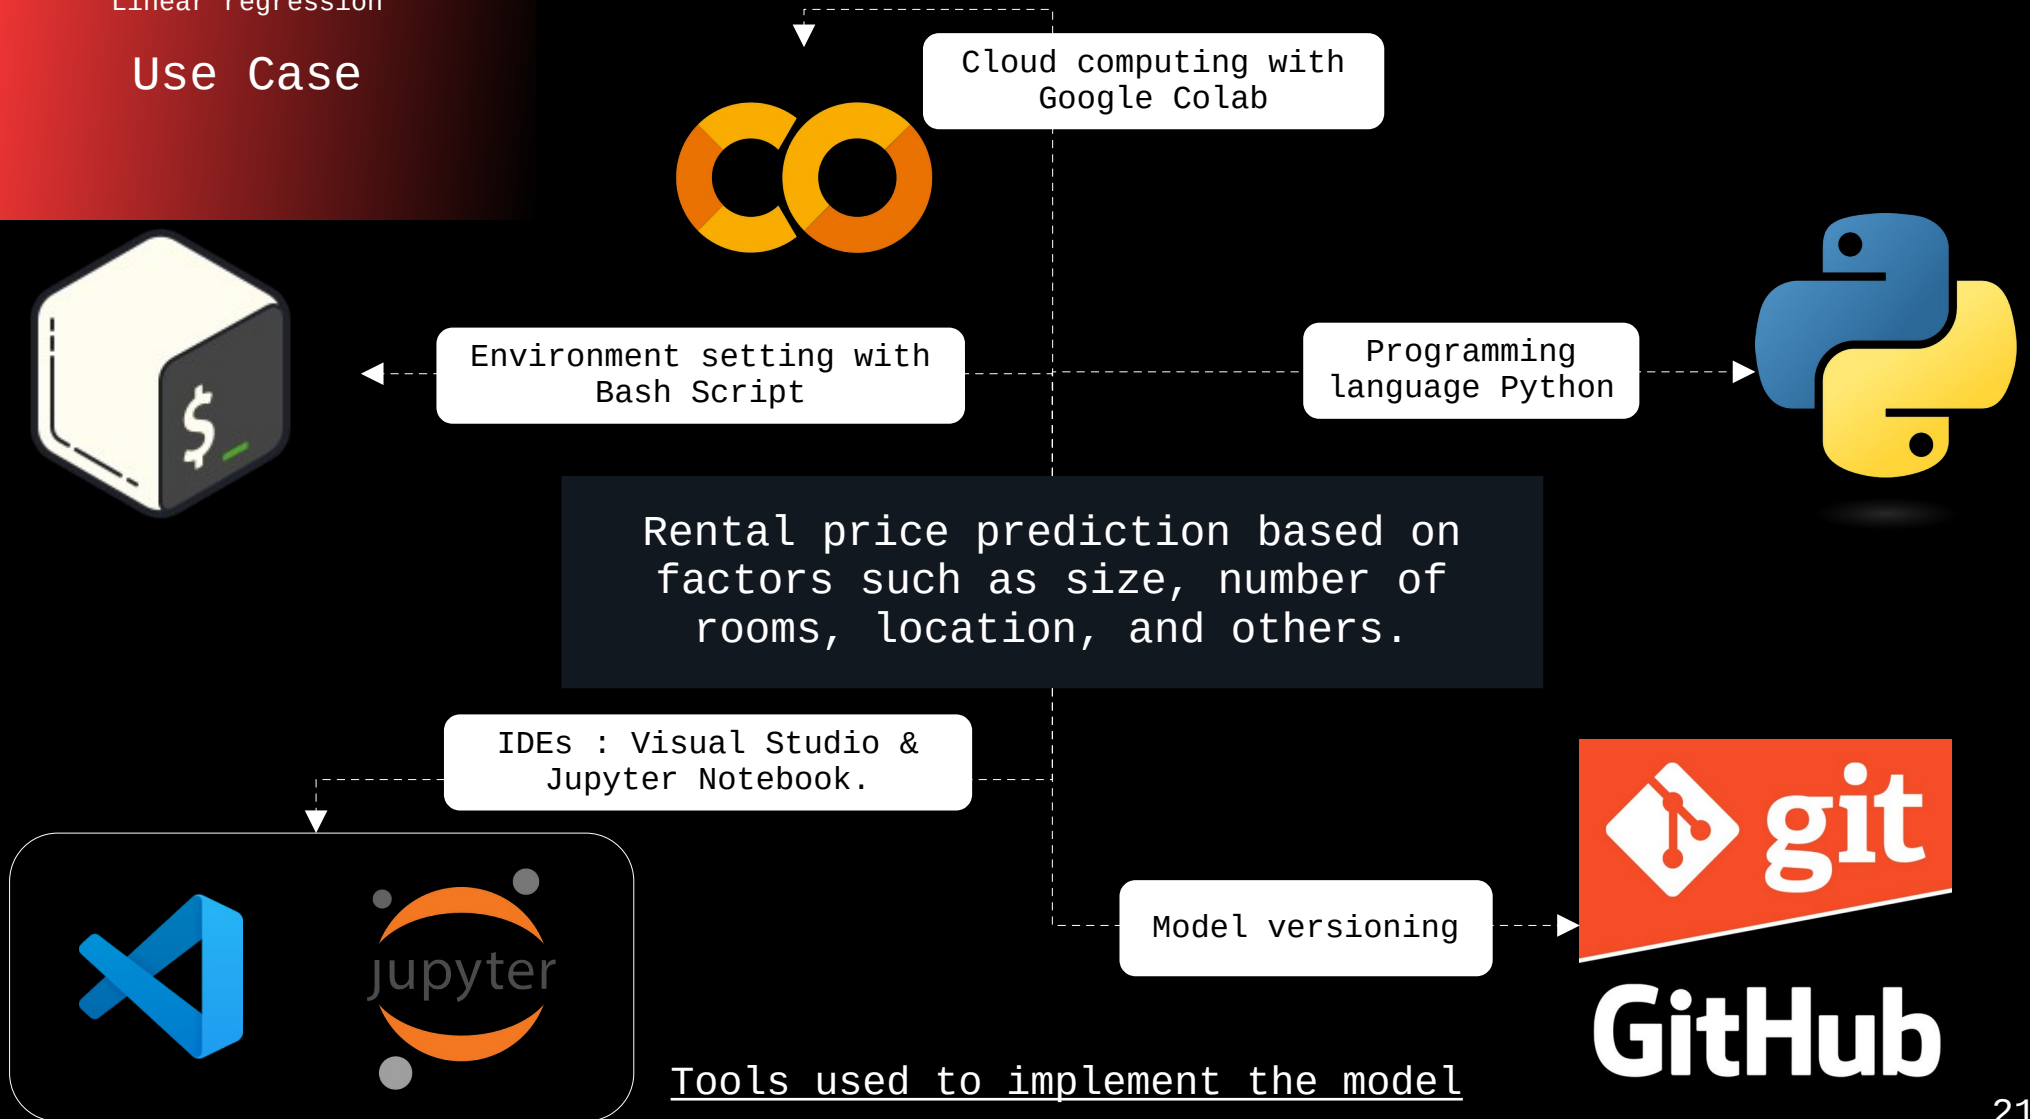
$$X = [x_1, x_2, \dots, x_k] \rightarrow X^T = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \end{bmatrix}$$

$$\alpha_i = [y_i - (X_i \theta + \hat{b})]$$

$$\nabla_{\theta} E(\theta) = \frac{-2}{I} \sum_{i=1}^I [y_i - (X_i \theta + \hat{b})] (X_i)^T = \frac{-2}{I} \sum_{i=1}^I \alpha_i \begin{bmatrix} x_1^i \\ x_2^i \\ \vdots \\ x_k^i \end{bmatrix}$$

$$\theta = \theta - \ni \nabla_{\theta} E(\theta) = \theta - \ni \frac{-2}{I} \sum_{i=1}^I \alpha_i \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \end{bmatrix} = \begin{bmatrix} \hat{a}_1 \\ \hat{a}_2 \\ \vdots \\ \hat{a}_k \end{bmatrix} - \ni \frac{-2}{I} \sum_{i=1}^I \alpha_i \begin{bmatrix} x_1^i \\ x_2^i \\ \vdots \\ x_k^i \end{bmatrix}_{20}$$

Use Case



- *Python is a programming language coded in C;*
- *Imperative and interpreted language;*
- *Most used language in data science;*
- *Includes various libraries, such as Pandas for structured data manipulation and NumPy for mathematical calculations, among others.*

Use Case

```

observations=[ 'pieds_carres','nombre_chambres','nombre_etages','distance_centre_ville','annee_construction','taille_garage','score_localisation','avec_piscine']
# observations=[ 'pieds_carres','nombre_etages','distance_centre_ville','annee_construction']

X=data[observations].values#Observations

print('Exemple des observations/features avant la standardization.')
print(X[0])

scaler = StandardScaler()
X_ = scaler.fit_transform(X)
print('\n')
print('Exemple des observations/features après la standardization.')
print(X_[0])

y=data[['prix']].values#Target (cible)
#2% des données (soit 100 échantillons/samples) seront utilisées pour le test.
#random_state permet de reproduire la distribution des données d'entraînement et de test.
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=1234)

X_train=scaler.fit_transform(X_train)#Standardiser les données d'entraînement.

```

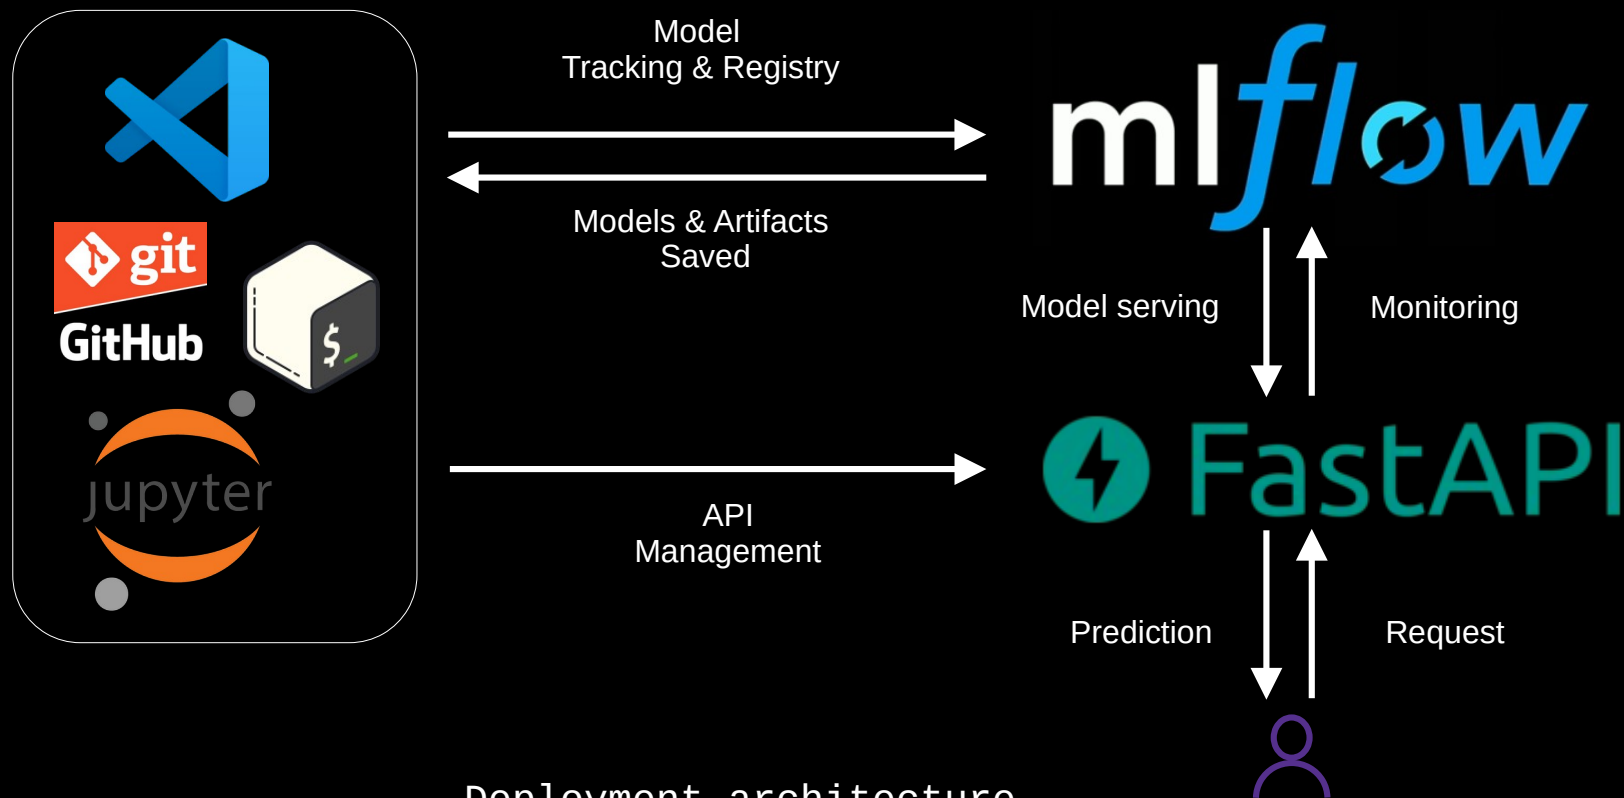
Let's code

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import mlflow
import numpy as np
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from datetime import time
from sklearn.preprocessing import StandardScaler
import random

```

Model deployment



Deployment architecture

